


Deep Learning (for Computer Vision)

Arjun Jain

Computer Vision: Working with Images

Picture Elements - PIXEL



PIXELS are ATOMIC ELEMENTS of an image.

In late 1960s, terminology 'pixel' was introduced by a group of scientists at JPL in California!

Image Types: Scalar and Binary

- A scalar image has $2^a - 1$ integer values

$$u \in \{0, 1, \dots, 2^a - 1\}$$

- a: level (bit)
- **Ex.** If 8 bit (a=8), image spans from 0 to 255
 - **0 black**
 - **255 white**
- **Ex.** If 1 bit (a=1), it is binary image, 0 and 1 only

Image Type: RGB (red, green, blue)



Image has three channels (bands), each channel spans a-bit values

Image format

- Some formats: TIF, PGM, PBM, GIF, JPEG, PNG, RAW etc.
- Medical Images: DICOM, Analyze, NIFTI etc.
- **HEADER:** contains image information, image size, pixel size, ...
- **DATA:** integer, double, float, unsigned integer, char,...

Practice: Image Format/Read/Show

```
from scipy import misc
l = misc.lena()
misc.imsave( 'lena.png', l) #uses the image module (PIL)

import matplotlib.pyplot as plt
plt.imshow(l)
plt.show()
```



PIL: Python Imaging Library

```
from PIL import Image
Img = Image.open('empire.jpg')
```



Matplotlib is a good graphics library with much
More powerful features than the

Pictures


- Computers use discrete form of the pictures
- The process transforming **continuous space** into **discrete space** is called **digitization**



Pictures



Digitization of an arc





Definition

- A (2D) picture P is a function defined on a (finite) rectangular subset G of a regular planar orthogonal array. G is called (2D) **grid**, and **an element of G is called pixel**. P assigns a value of $P(p)$ to each point



Definition

- A (2D) picture P is a function defined on a (finite) rectangular subset G of a regular planar orthogonal array. G is called (2D) **grid**, and **an element of G is called pixel**. P assigns a value of $P(p)$ to each point





Definition

- Pictures are not only sampled, they are also quantized: they may have only a finite number of possible values (i.e., 0 to 255, 0-1, ...)




Resolution

- Resolution is a display parameter, defined in dots per inch (DPI) or equivalent measures of spatial pixel density, and its standard value for recent screen technologies is 72 dpi. Recent printer resolutions are in 300 dpi and/or 600 dpi.




Filtering

- An image contains discrete number of pixels
 - A simple example
 - Pixel value:
 - “grayscale”
- (or “intensity”): [0,255]




Filtering

- An image contains discrete number of pixels
 - A simple example
 - Pixel value:
 - “grayscale”
(or “intensity”): [0,255]
 - “color”
 - RGB: [R, G, B]
 - Lab: [L, a, b]
 - HSV: [H, S, V]



Filtering



Filtering : RGB Channels




Image Histogram



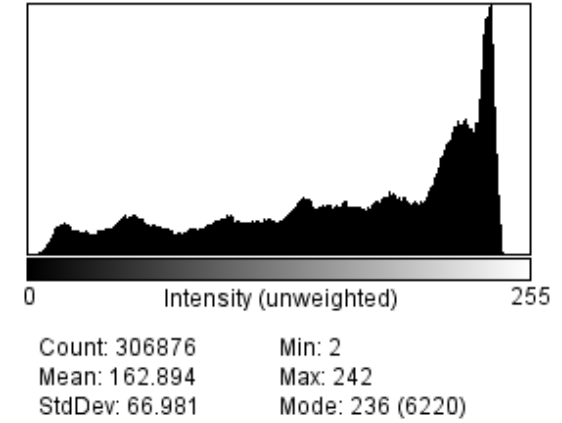



Image Histogram

Use ImageJ and/or FIJI



Correlation

$$I \otimes W = \sum_k \sum_l I(k, l) W(i + k, j + l)$$

I = Image

W = Kernel

I

i_1	i_2	i_3
i_4	i_5	i_6
i_7	i_8	i_9

\otimes

W

w_1	w_2	w_3
w_4	w_5	w_6
w_7	w_8	w_9

\longrightarrow

$$\begin{aligned}
 I * W = & i_1 w_1 + i_2 w_2 + i_3 w_3 \\
 & + i_4 w_4 + i_5 w_5 + i_6 w_6 \\
 & + i_7 w_7 + i_8 w_8 + i_9 w_9
 \end{aligned}$$

Convolution

$$I \otimes W = \sum_k \sum_l I(k, l) W(i - k, j - l)$$

I = Image

W = Kernel

I

f_1	f_2	f_3
f_4	f_5	f_6
f_7	f_8	f_9



w_9	w_8	w_7
w_6	w_5	w_4
w_3	w_2	w_1

\swarrow Y - flip

w_7	w_8	w_9
w_4	w_5	w_6
w_1	w_2	w_3

\swarrow X - flip



w_1	w_2	w_3
w_4	w_5	w_6
w_7	w_8	w_9

w

$$\begin{aligned} I * W = & i_1 w_9 + i_2 w_8 + i_3 w_7 \\ & + i_4 w_6 + i_5 w_5 + i_6 w_4 \\ & + i_7 w_3 + i_8 w_2 + i_9 w_1 \end{aligned}$$

Convolution

$$I(x, y) * W = I(x + 1, y + 1)W(-1, -1) + I(x, y + 1)W(0, -1) + I(x - 1, y + 1)W(1, -1) + \\ I(x + 1, y)W(-1, 0) + I(x, y)W(0, 0) + I(x - 1, y)W(1, 0) + \\ I(x + 1, y - 1)W(-1, 1) + I(x, y - 1)W(0, 1) + I(x - 1, y - 1)W(1, 1)$$



$$I * W = \sum_{i=-1}^1 \sum_{j=-1}^1 I(x - i, y - i)W(i, j)$$


Coordinates

-1,0	0,1	1,1
-1,0	0,0	1,0
-1,-1	0,-1	1,-1

Filtering

Filtering

- Modify pixels based on some function of the neighbourhood



Filtering

- The output is the linear combination of the neighbourhood pixels

1	3	0
2	10	2
4	1	1

\otimes

1	0	-1
1	0.1	-1
1	0	-1


$=$

Image

Kernel

Filtering


- The output is the linear combination of the neighbourhood pixels



Filtering examples



*




0	0	0
0	1	0
0	0	0

=

Filtering examples



*



0	0	0
0	1	0
0	0	0


=



Filtering examples



*




0	0	0
1	0	0
0	0	0

=

Filtering examples



*




0	0	0
1	0	0
0	0	0

=



Filtering examples



$\ast \frac{1}{9}$

1	1	1
1	1	1
1	1	1

=


Filtering examples




$\ast \frac{1}{9}$

1	1	1
1	1	1
1	1	1

=



Filtering examples



$$* \frac{1}{25} =$$

1	1	1
1	1	1
1	1	1

Filtering examples




$$\begin{matrix}
 \text{3D volume of 25 colored blocks (5x5x1)} \\
 * \frac{1}{25} \\
 \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}
 \end{matrix}
 =$$



Filtering examples - Gaussian



*



=



Filtering example – Gaussian vs. Smoothing



Gaussian Smoothing



Smoothing by Averaging

Filtering example – Noise filtering



Gaussian Smoothing



Smoothing by Averaging

Filtering example – Noise filtering



Gaussian Noise



After averaging



After Gaussian Smoothing

Thank you!